

# The Universe of Possibilities Just Widened with Machine Learning

Using Machine Learning to make the world a better place and flesh out our understanding of ourselves

Please sign-in: [ucfai.org/signin](https://ucfai.org/signin)

By: John Muchovej

# The Universe of Possibilities Just Widened with Machine Learning

## Part 1: Intro the Machine Learning (ML)

Please sign-in: [ucfai.org/signin](https://ucfai.org/signin)

# Who am I and why should you trust me?

- John Muchovej
  - Comp Sci & Math-Econ major
- Machine Learning Researcher (2 labs)
  - Computer Vision
  - Natural Language Processing
- Freelance Data Scientist
- Contractor for Udacity, mentoring 300+ students in 10 “Nanodegrees”
  - namely Machine Learning, Deep Reinforcement Learning, and the like
- Founder of AI@UCF, sponsored by CBMM@MIT and Intel

**Quick note:**

*A single* deck will  
be used for both  
Part 1 and Part 2

# Disclaimer:

An information dense deck lies ahead, **don't let that discourage you**. I'm giving you **words & resources** to **learn about after KH3**.

**At the end of the slide deck,  
you'll find links to helpful  
resources to learn more  
about this on your own!**

# Part 1: Intro to Machine Learning (ML)

# Agenda

- What is Machine Learning? (ML)
- What can you use it for?
- Downloading datasets
- Exploring data
- What are Neural Networks?
- Building Neural Networks in PyTorch



# What is Machine Learning?

- A toolset which avoids explicit programming by **learning from data**
- Modern machine learning is **approximates functions**
- Machine learning is much less about your algorithm of choice and more about ***the data you have access to***

More data beats a cleverer algorithm [1]

- Pedro Domingos, 2012

# What can we use ML for?

## Computer Vision

- Object detection
- Object recognition
- Object segmentation
- Object classification
- Landmark detection
- Optical Character Recognition
- ...


## Natural Language Processing

- Sentiment Analysis
- Named Entity Recognition
- Syntax Analysis
- Content classification
- ...

**There's significantly more which you can do, beyond what's enumerated here.**

**As long as you have data**, you can probably use ML for it! :D

# The Process (Supervised Learning)



# Downloading Datasets

- Finding data can be somewhat difficult
- Two great resources, though:
  - a. Kaggle ([kaggle.com/datasets](https://kaggle.com/datasets))
  - b. Awesome Public Datasets, GitHub ([awesomedata/awesome-public-datasets](https://github.com/awesomedata/awesome-public-datasets))
- Generally, data can come in all kinds of forms
  - a. For a hackathon, though, try to snag tabular or image data
  - b. Audio can be pretty cumbersome if this is your first exposure to ML

# Let's pop over to Jupyter Notebooks for “Exploring data”

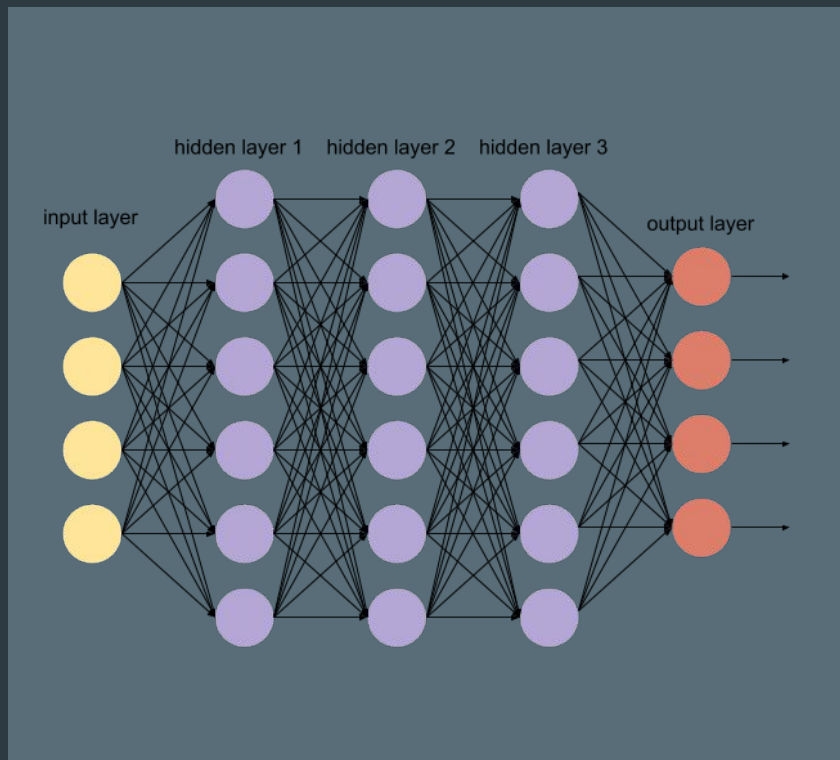
# What are Neural Networks? (NNs)

- They've turned ML from your standard research field into the benchmarks-land-on-the-front-page-field we have today
- They have incredible representation power
  - More on this in a bit
- They can be increasingly sped by using GPUs
- One of many, many ML, “Learning Algorithms”
  - Support Vector Machines
  - k-Nearest Neighbors

**Suffice it to say, NNs are just a fraction of what *is* considered “ML.”**

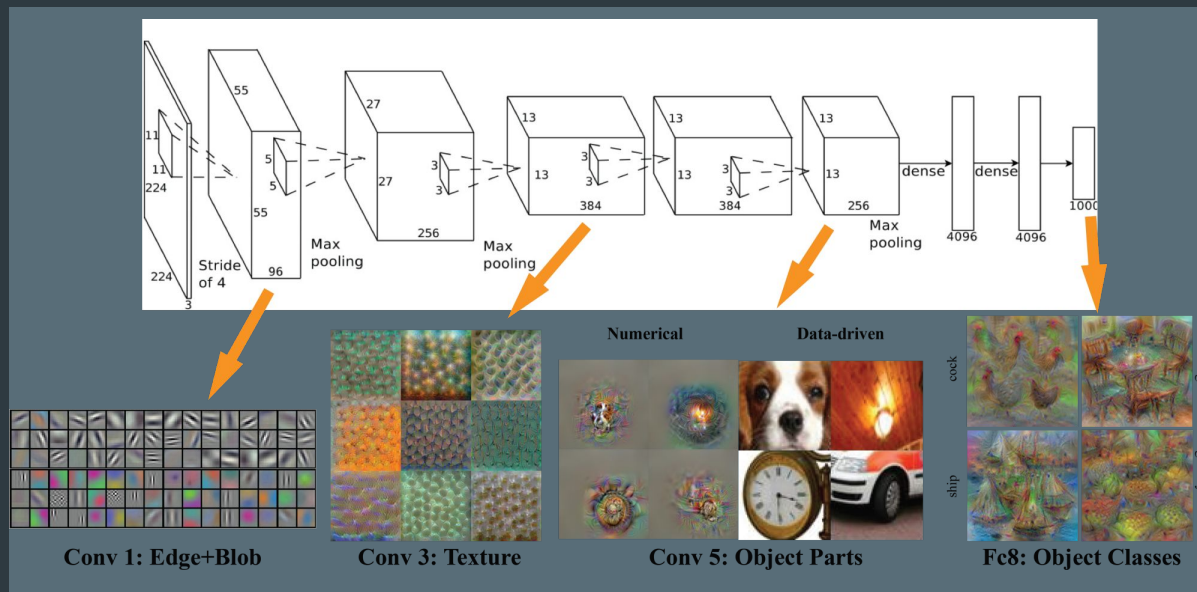
# More technically, Neural Networks (NNs / FCNs)

- NNs take in data, as an array (or vector)
- They multiply and combine the data you give it to produce new “features”
- The weights are “learned” through optimization
  - This flavor of optimization is called “Gradient Descent”



# Convolutional Neural Networks (CNNs)

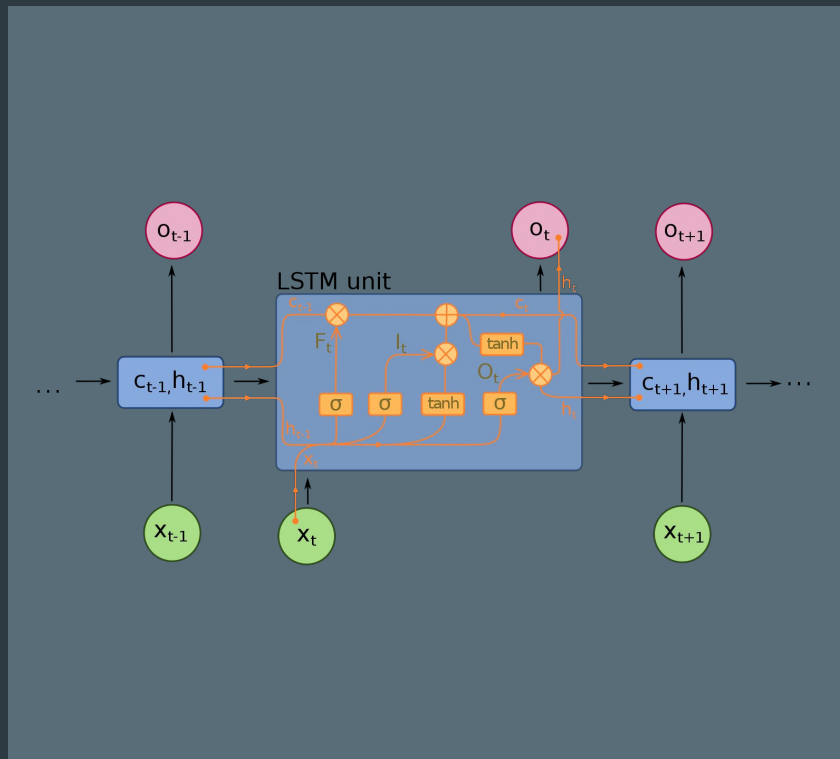
- NNs remove spatial data by flattening their input
- CNNs overcome that
- Typically called “feature extractors”
- Goal of CNNs is to extract high-level objects to be used by the FCNs





# Recurrent Neural Networks (RNNs)

- Neither CNNs nor FCNs can handle sequences
- RNNs overcome that by maintaining “state”
- The most commonly used RNNs:
  - LSTM (Long Short-Term Memory)
  - GRU (Gated Recurrent Unit)



# The nitty-gritty (i.e. hyperparameters)

- Typical optimization algorithm is Adam [2]
- Typical activation function for non-output nodes is ReLU [3]
- Loss functions: ([PyTorch documentation](#))
  - Loss functions are used to actually train your model
  - When doing classification, consider using...
    - *Cross-Entropy*
  - When doing regression, consider using...
    - *Mean-Squared Error (MSE)*
    - *L1*
    - *SmoothL1*
      - From personal experience, this one seems to do best

# Let's familiarize ourselves with the PyTorch Framework

**Let us know what you thought!**  
**[ucfai.org/feedback](https://ucfai.org/feedback)**

# The Universe of Possibilities Just Widened with Machine Learning

## Part 2: DevOps for Machine Learning

Please sign-in: [ucfai.org/signin](https://ucfai.org/signin)

# Who am I and why should you trust me?

- John Muchovej
  - Comp Sci & Math-Econ major
- Machine Learning Researcher (2 labs)
  - Computer Vision
  - Natural Language Processing
- Freelance Data Scientist
- Contractor for Udacity, mentoring 300+ students in 10 “Nanodegrees”
  - namely Machine Learning, Deep Reinforcement Learning, and the like
- Founder of AI@UCF, sponsored by CBMM@MIT and Intel

## Part 2: DevOps for ML

**Quick note:**

This is Part 2, if you  
want to learn ML,  
checkout Part 1



# Agenda

- DevOps is the bottleneck of ML
- Setting up a Google Cloud Platform account (GCP)
- Selecting resources for GCE (Compute Engine)
- Using GPUs in PyTorch
- High-level exploration of the different GCP APIs

# DevOps is the bottleneck of ML

- Getting your models up and running can be a pain on the cloud
- Ensuring your development and production environment are as similar as possible is challenging
- Check-out open-source “cutting edge research” – try reproducing their results and running their code
  - Actually, don't. It's a pain, liable to cause hemorrhaging

# Setting up a Google Cloud Platform account (GCP)

- Those Anaconda YAML files from earlier should be enough to get you started on GCP!
  - I've included a setup script which will also get you running with PyTorch, CUDA, and the like. :D
- Launch the "Compute Engine"
- Go to your Quotas and request...
  - "Global GPU limit" be raised to 3-4
  - "us-\*#-(a/b/...)" zones raised to 2-3 (take a look at [GPU locations](#))

# Selecting resources for GCE (Compute Engine)

- Start out with 4-cores, 16GB RAM (n1-standard-4)
- Attach a single NVIDIA T4 / P4 (this will run you \$0.79-1.14 / hour)
- **Focus on deploying multiple experiments with different “hyperparameters”**

[Google Cloud Compute Engine pricing / etc.](#)

- [GPU pricing](#)
- [CPU / RAM pricing](#)

# Using GPUs in PyTorch

- The [PyTorch documentation](#) on this is great
- General rules of thumb, though
  - Put your model on the GPU
  - Process your data on CPU, then transfer it to the GPU
  - In a time-crunch like this, stick to 1 GPU

# High-level exploration of the different GCP APIs

If you're doing a project, like an Alexa skill, odds are the GCP ML APIs should do enough!

- [AutoML](#) is a really interesting tool
  - It allows you to perform “transfer learning” by uploading your own data to optimize Google’s pre-existing models
  - [AutoML Vision \(Beta\)](#)
  - [AutoML Natural Language \(Beta\)](#)
  - [AutoML Translation \(Beta\)](#)
- As a researcher, I haven’t really used these, however I’ve read/heard that the Python library is okay and than raw REST requests might be more useful

# Questions?



**Let us know what you thought!**  
**[ucfai.org/feedback](https://ucfai.org/feedback)**